

```
import tensorflow as tf

# Display the version
print(tf.__version__)

# other imports
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.models import Model
```

2.19.0

```
# Load in the data
cifar10 = tf.keras.datasets.cifar10

# Distribute it to train and test set
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
print(x_train.shape, y_train.shape, x_test.shape, y_test.shape)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 ————— 4s 0us/step
(50000, 32, 32, 3) (50000, 1) (10000, 32, 32, 3) (10000, 1)

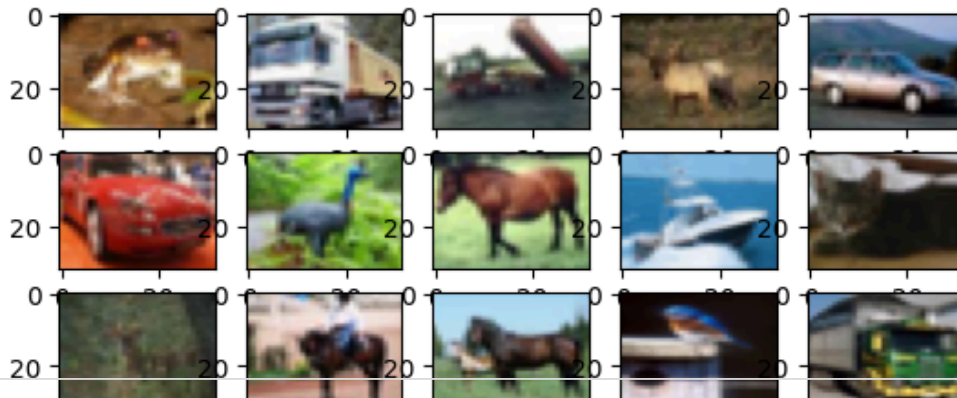
```
# Reduce pixel values
x_train, x_test = x_train / 255.0, x_test / 255.0

# flatten the label values
y_train, y_test = y_train.flatten(), y_test.flatten()
```

```
# visualize data by plotting images
fig, ax = plt.subplots(5, 5)
k = 0

for i in range(5):
    for j in range(5):
        ax[i][j].imshow(x_train[k], aspect='auto')
        k += 1

plt.show()
```



```
# number of classes
K = len(set(y_train))

# calculate total number of classes
# for output layer
print("number of classes:", K)

# Build the model using the functional API
# input layer
i = Input(shape=x_train[0].shape)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(i)
x = BatchNormalization()(x)
x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2))(x)

x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2))(x)

x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = Conv2D(128, (3, 3), activation='relu', padding='same')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2))(x)

x = Flatten()(x)
x = Dropout(0.2)(x)

# Hidden layer
x = Dense(1024, activation='relu')(x)
x = Dropout(0.2)(x)

# last hidden layer i.e.. output layer
x = Dense(K, activation='softmax')(x)

model = Model(i, x)

# model description
model.summary()
```

number of classes: 10
Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None , 32, 32, 3)	0
conv2d (Conv2D)	(None , 32, 32, 32)	896
batch_normalization (BatchNormalization)	(None , 32, 32, 32)	128
conv2d_1 (Conv2D)	(None , 32, 32, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None , 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None , 16, 16, 32)	0
conv2d_2 (Conv2D)	(None , 16, 16, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None , 16, 16, 64)	256
conv2d_3 (Conv2D)	(None , 16, 16, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None , 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None , 8, 8, 64)	0
conv2d_4 (Conv2D)	(None , 8, 8, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None , 8, 8, 128)	512
conv2d_5 (Conv2D)	(None , 8, 8, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None , 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None , 4, 4, 128)	0
flatten (Flatten)	(None , 2048)	0
dropout (Dropout)	(None , 2048)	0
dense (Dense)	(None , 1024)	2,098,176
dropout_1 (Dropout)	(None , 1024)	0
dense_1 (Dense)	(None , 10)	10,250

```
# Compile
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
# Fit
r = model.fit(
    x_train, y_train, validation_data=(x_test, y_test), epochs=5)
```

```
Epoch 1/5
1563/1563 ————— 411s 263ms/step - accuracy: 0.7802 - loss: 0.6436 - val_accu
Epoch 2/5
1563/1563 ————— 409s 261ms/step - accuracy: 0.8156 - loss: 0.5345 - val_accu
Epoch 3/5
1563/1563 ————— 410s 262ms/step - accuracy: 0.8424 - loss: 0.4505 - val_accu
Epoch 4/5
1563/1563 ————— 443s 263ms/step - accuracy: 0.8695 - loss: 0.3730 - val_accu
Epoch 5/5
1563/1563 ————— 442s 263ms/step - accuracy: 0.8910 - loss: 0.3133 - val_accu
```

```
# Fit with data augmentation
# Note: if you run this AFTER calling
# the previous model.fit()
# it will CONTINUE training where it left off
batch_size = 32
data_generator = tf.keras.preprocessing.image.ImageDataGenerator(
    width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)


train_generator = data_generator.flow(x_train, y_train, batch_size)
steps_per_epoch = x_train.shape[0] // batch_size

r = model.fit(train_generator, validation_data=(x_test, y_test),
              steps_per_epoch=steps_per_epoch, epochs=5)
```

```
Epoch 1/5
/usr/local/lib/python3.12/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.
self._warn_if_super_not_called()
1562/1562 ————— 454s 291ms/step - accuracy: 0.7795 - loss: 0.6734 - val_accu
Epoch 2/5
1/1562 ————— 5:50 225ms/step - accuracy: 0.8438 - loss: 0.4647/usr/local/l
self._interrupted_warning()
1562/1562 ————— 19s 12ms/step - accuracy: 0.8438 - loss: 0.4647 - val_accu
Epoch 3/5
1562/1562 ————— 431s 276ms/step - accuracy: 0.8067 - loss: 0.5765 - val_accu
Epoch 4/5
1562/1562 ————— 19s 12ms/step - accuracy: 0.9062 - loss: 0.3415 - val_accu
Epoch 5/5
1562/1562 ————— 428s 274ms/step - accuracy: 0.8170 - loss: 0.5398 - val_accu
```

```
# Plot accuracy per iteration
plt.plot(r.history['accuracy'], label='acc', color='red')
plt.plot(r.history['val_accuracy'], label='val_acc', color='green')
plt.legend()
```

<matplotlib.legend.Legend at 0x789f01ccf890>



```
# label mapping

labels = '''airplane automobile bird cat deerdog frog horseship truck'''.split()

# select the image from our test dataset
image_number = 0

# display the image
plt.imshow(x_test[image_number])

# load the image in an array
n = np.array(x_test[image_number])


# reshape it
p = n.reshape(1, 32, 32, 3)

# pass in the network for prediction and
# save the predicted label
predicted_label = labels[model.predict(p).argmax()]

# load the original label
original_label = labels[y_test[image_number]]


# display the result
print("Original label is {} and predicted label is {}".format(
    original_label, predicted_label))
```

1/1 ————— 0s 253ms/step
Original label is cat and predicted label is frog



```
# save the model  
model.save('geeksforgeeks.h5')
```

20 +
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.s



```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import tensorflow as tf
from tensorflow import keras
import seaborn as sns
import os
from datetime import datetime

import warnings
warnings.filterwarnings("ignore")
```

```
import pandas as pd

data = pd.read_csv('/content/all_stocks_5yr.csv', delimiter=',', on_bad_lines='skip')
print(data.shape)
print(data.sample(7))
```

```
(152910, 7)
```

	date	open	high	low	close	volume	Name
50493	2014-06-02	108.56	108.95	107.8100	108.35	1270180	ANTM
74656	2015-03-16	67.76	68.48	67.7500	68.43	2966222	BAX
149681	2016-09-22	27.08	27.72	26.9500	27.61	5094979	CTL
72524	2016-09-23	46.43	46.87	46.3100	46.61	1957032	A
27624	2017-10-23	100.04	100.04	99.1200	99.49	463664	AIZ
42411	2017-04-26	165.61	165.61	164.3397	164.61	4175536	AMGN
76393	2017-02-06	162.42	164.08	162.3800	163.98	3110494	BA

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 152910 entries, 0 to 152909
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  -
0   date    152910 non-null     object
1   open    152909 non-null     float64
2   high    152909 non-null     float64
3   low     152909 non-null     float64
4   close   152910 non-null     float64
5   volume  152910 non-null     int64
6   Name    152909 non-null     object
dtypes: float64(4), int64(1), object(2)
memory usage: 8.2+ MB
```

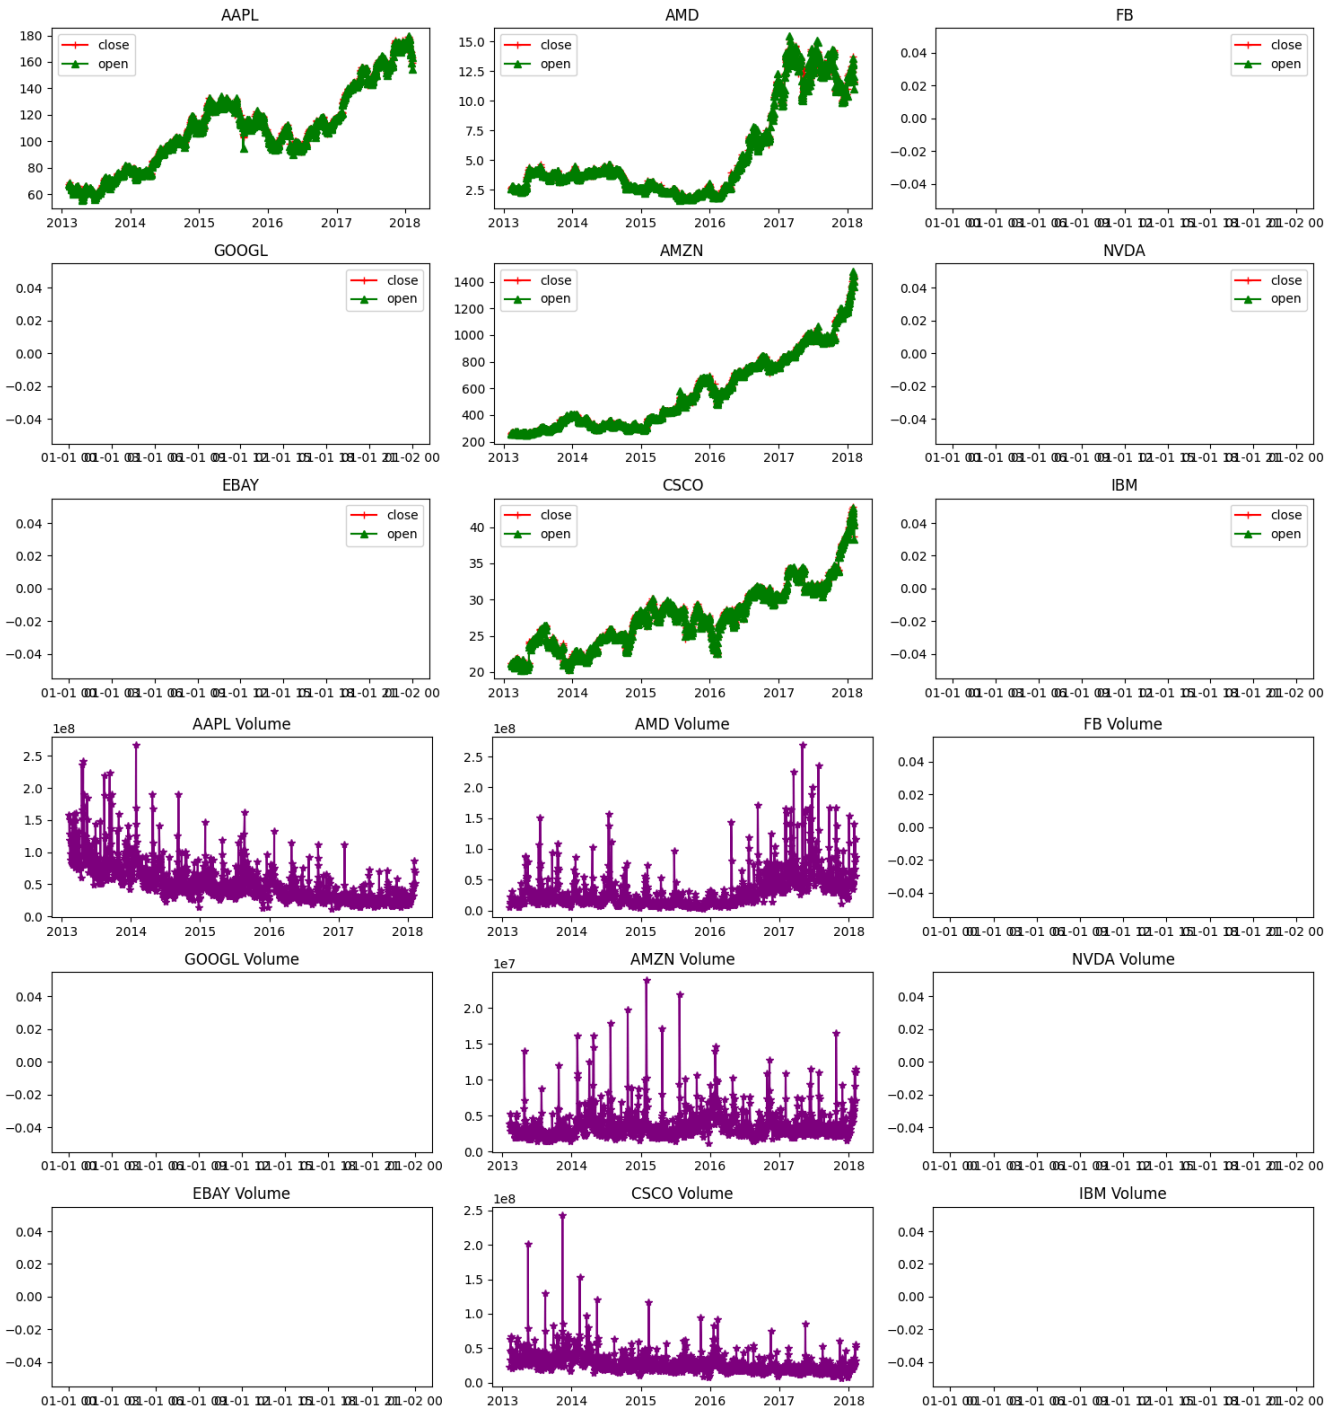
```
data['date'] = pd.to_datetime(data['date'])
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 152910 entries, 0 to 152909
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  -
0   date    152910 non-null     datetime64[ns]
1   open    152909 non-null     float64
2   high    152909 non-null     float64
3   low     152909 non-null     float64
4   close   152910 non-null     float64
5   volume  152910 non-null     int64
6   Name    152909 non-null     object
dtypes: datetime64[ns](1), float64(4), int64(1), object(1)
memory usage: 8.2+ MB
```


```
companies = ['AAPL', 'AMD', 'FB', 'GOOGL', 'AMZN', 'NVDA', 'EBAY', 'CSCO', 'IBM']
```

```
plt.figure(figsize=(15, 8))
for index, company in enumerate(companies, 1):
    plt.subplot(3, 3, index)
    c = data[data['Name'] == company]
    plt.plot(c['date'], c['close'], c="r", label="close", marker="+")
    plt.plot(c['date'], c['open'], c="g", label="open", marker="^")
    plt.title(company)
    plt.legend()
    plt.tight_layout()
```


```
plt.figure(figsize=(15, 8))
for index, company in enumerate(companies, 1):
    plt.subplot(3, 3, index)
    c = data[data['Name'] == company]
    plt.plot(c['date'], c['volume'], c='purple', marker='*')
    plt.title(f"{company} Volume")
    plt.tight_layout()
```



```
plt.figure(figsize=(15, 8))
for index, company in enumerate(companies, 1):
    plt.subplot(3, 3, index)
    c = data[data['Name'] == company]
    plt.plot(c['date'], c['volume'], c='purple', marker='*')
    plt.title(f"{company} Volume")
plt.tight_layout()
```

```
apple = data[data['Name'] == 'AAPL']
prediction_range = apple.loc[(apple['date'] > datetime(2013,1,1))
                             & (apple['date'] < datetime(2018,1,1))]
plt.plot(apple['date'],apple['close'])
plt.xlabel("Date")
plt.ylabel("Close")
plt.title("Apple Stock Prices")
plt.show()
```



```
close_data = apple.filter(['close'])
dataset = close_data.values
training = int(np.ceil(len(dataset) * .95))
print(training)
```

1197

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(dataset)

train_data = scaled_data[0:int(training), :]
# prepare feature and labels
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])

x_train, y_train = np.array(x_train), np.array(y_train)
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))
```

```
model = keras.models.Sequential()
model.add(keras.layers.LSTM(units=64,
                             return_sequences=True,
```

```

        return_sequences=True,
        input_shape=(x_train.shape[1], 1)))
model.add(keras.layers.LSTM(units=64))
model.add(keras.layers.Dense(32))
model.add(keras.layers.Dropout(0.5))
model.add(keras.layers.Dense(1))
model.summary

```

```

keras.src.models.model.Model.summary
def summary(line_length=None, positions=None, print_fn=None, expand_nested=False,
            show_trainable=False, layer_range=None)

```

[/usr/local/lib/python3.12/dist-packages/keras/src/models/model.py](#)
Prints a string summary of the network.

Args:

- line_length: Total length of printed lines (e.g. set this to adapt the display to different

```

model.compile(optimizer='adam',
              loss='mean_squared_error')
history = model.fit(x_train,
                    y_train,
                    epochs=10)

```

```

Epoch 1/10
36/36 ————— 6s 58ms/step - loss: 0.0533
Epoch 2/10
36/36 ————— 2s 60ms/step - loss: 0.0087
Epoch 3/10
36/36 ————— 3s 92ms/step - loss: 0.0090
Epoch 4/10
36/36 ————— 4s 99ms/step - loss: 0.0090
Epoch 5/10
36/36 ————— 4s 58ms/step - loss: 0.0083
Epoch 6/10
36/36 ————— 3s 59ms/step - loss: 0.0075
Epoch 7/10
36/36 ————— 3s 70ms/step - loss: 0.0071
Epoch 8/10
36/36 ————— 3s 85ms/step - loss: 0.0065
Epoch 9/10
36/36 ————— 2s 59ms/step - loss: 0.0061
Epoch 10/10
36/36 ————— 2s 59ms/step - loss: 0.0066

```

```

test_data = scaled_data[training - 60:, :]
x_test = []
y_test = dataset[training:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

x_test = np.array(x_test)
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

mse = np.mean(((predictions - y_test) ** 2))
rmse = np.sqrt(mse)

print("MSE", mse)
print("RMSE", np.sqrt(mse))

```

```

2/2 ————— 1s 425ms/step
MSE 45.63328004905825
RMSE 6.755240931977056

```

```

train = apple[:training]
test = apple[training:]
test['Predictions'] = predictions

plt.figure(figsize=(10, 8))
plt.plot(train['date'], train['close'])
plt.plot(test['date'], test[['close', 'Predictions']])
plt.title('Apple Stock Close Price')
plt.xlabel('Date')
plt.ylabel('Close')
plt.legend(['Train', 'Test', 'Predictions'])

```

<matplotlib.legend.Legend at 0x7dc867a727e0>

Apple Stock Close Price

